

title: "GPTs Unplugged: The Wild Ride of AI Chatbots (and How OverKill Hill Is Pushing the Limits)"
description: "A punchy, comprehensive journey through the evolution of Generative Pre-trained Transformers (GPTs) – what they are, how they rose to fame, why they matter, and how they stack up against Microsoft's Copilot Studio Lite and Google's Gemini Gems. Includes a spotlight on OverKill Hill P³'s unique contributions (Gleefully Tools, BFS, BrandGuard, Foundry, Ask Jamie, HomeEstimate-R, Unnocked Truth) and the new era of custom AI assistants." author: "OverKill Hill P³" keywords: "GPT, Generative AI, ChatGPT, Google Gemini, Copilot Studio, OverKill Hill, Gleefully Tools, AI Evolution, AI Chatbots, Custom GPTs"

Introduction: From Nerdy Novelty to Global Phenomenon

Figure 1: OverKill Hill P³ banner – where AI architecture meets attitude. OverKill Hill's mission is to bring powerful, scalable AI tools to those seeking smarter systems ¹.

Once upon a time (okay, **2018**), "GPT" was just another obscure acronym floating around AI labs. Fast forward to today and GPTs are **everywhere** – writing your emails, acing law exams, even helping code the next big app. How the heck did we get here? In late 2022, a chatbot called **ChatGPT** seemingly came out of nowhere and took the world by storm. It hit 100 million users in two months flat ² – officially the fastest-growing consumer app in history, outpacing the likes of TikTok and Instagram. Suddenly, everyone and their grandma heard about AI that *talks back*.

This paper is your irreverent guide to that whirlwind. We'll break down **what GPTs are** (in plain English, with only minimal brain-bleed tech jargon), **how they evolved** from a quirky research idea to the engines inside today's most powerful chatbots, and **why they matter** beyond the tech hype. Along the way, we'll pit GPTs against the other new AI sidekicks on the block – namely Microsoft's **Copilot Studio Lite** and Google's **Gemini Gems** – and see how they compare. Spoiler: OpenAI's GPT-based ChatGPT and Google's Gem-powered assistants are currently the twin titans ruling the public AI playground, while Microsoft's approach is a bit more buttoned-up for the enterprise crowd. We'll explore why that is.

And because this journey wouldn't be complete without some home-grown flavor, we're highlighting how **OverKill Hill P³** (yours truly) is innovating in the GPT space. From the whimsical productivity pal **Gleefully Tools** to the no-nonsense **BrandGuard** content watchdog, from our future-focused **BFS** ("Framing Intelligent Futures") initiative to custom assistants like **Ask Jamie**, **HomeEstimate-R**, and the cheekily named **Unnocked Truth**, we'll show how a little irreverence and a lot of creativity can push GPT tech into new realms. These aren't just theoretical – they're real GPT-powered tools (hosted on chatgpt.com as custom GPTs) that demonstrate what's possible when you mix cutting-edge AI with a bold vision and a sense of humor.

Buckle up, because we're about to cover a lot of ground – roughly the last seven years of AI breakthroughs – in language anyone can follow. Think of it as chatting with a (somewhat hyper-caFFEinated) friend who's super into AI, rather than slogging through a dry textbook. We've also thrown in high-fidelity figures and

diagrams, so your eyeballs get a break from text now and then. By the end, you'll not only understand GPTs – the Generative Pre-trained Transformers – but also get why they're a huge deal and how they're changing the game.

Oh, and we'll keep it real: GPTs are awesome, but they're not magic or infallible. We'll discuss the *can-dos* and *can't-dos*, the wows and the uh-ohs (looking at you, AI hallucinations and bias).

Ready? Let's demystify these smarty-pants AIs and have some fun along the way.

What the Heck Is a GPT, Anyway?

So what exactly is a **GPT**? The acronym stands for **Generative Pre-trained Transformer**. Yeah, that's a mouthful. Let's unpack it in normal-people terms:

- **Generative:** It means this AI doesn't just regurgitate info – it **creates** stuff. Give it a prompt or question, and it will *generate* a fresh answer or piece of text for you. In other words, it's not picking from a fixed database of responses; it's making new sentences on the fly, like a super advanced predictive text engine.
- **Pre-trained:** Before you ever meet this model, it's been through **boot camp**. It spent ages gobbling up text from the internet, books, articles – basically half the internet – learning the statistical patterns of language. This “pre-training” is unsupervised, meaning it just learns from raw text without needing humans to label anything. By the end of it, the GPT has a general **feel** for how language works. Only after that do we fine-tune it (teach it specific tasks) with smaller, more focused datasets. The pre-training is the heavy lifting – it's why GPTs can do so many things out-of-the-box.
- **Transformer:** This is the name of the neural network architecture under the hood. The Transformer (introduced by Google researchers in 2017) was a game-changer in machine learning ³. Unlike older neural nets that processed words one by one in order (RNNs), transformers have this attention mechanism that lets them look at an entire sentence (or paragraph) at once and see how every word relates to the others ³. Think of it like having a million eyes scanning the text from all angles, rather than a single-track mind reading left-to-right. This architecture is **really good** at language – so good that OpenAI built GPTs on top of it, hence “Transformer” in the name.

In plainer terms, a GPT is like the world's most talented autocomplete. You give it some input (a prompt), and it tries to predict the most likely continuation or answer based on everything it “learned” from that massive corpus of text. But because it read so much and learned the *shape* of human language, those predictions can be eerily coherent – often indistinguishable from something a human might write.

Crucially, GPTs **aren't** copy-paste machines or fact databases. They don't have a directory of facts they pull from. They generate text by predicting next words, which means they can sometimes BS their way through an answer that *sounds* confident but is actually nonsense. (We'll talk about those “hallucinations” later – every magician has their flaws.) They also don't truly *understand* meaning the way humans do – they don't have beliefs or comprehension; they just have patterns. But with the right guidance, those patterns are incredibly useful.

It helps to visualize the Transformer brain inside a GPT. Here's a peek under the hood:

Figure 2: The Transformer neural network architecture underlying GPT models. It's basically a stack of encoder/decoder layers using multi-head self-attention mechanisms ⁴. In non-geek speak: the model is built in layers that all pay attention to different aspects of the input text, which allows the GPT to consider context from the entire input sequence at once, rather than forgetting things after a few words. This parallel processing of language was a huge leap over older RNN models ³.

When you type a sentence and ask a GPT to complete it, the Transformer is crunching through all the possible next words considering the whole context you gave. That's why, for example, if you start a prompt with "In a shocking turn of events, the President of the United States...", the GPT knows to expect something about politics or international events and not, say, a recipe for banana bread (unless the President *is* giving a banana bread recipe, which would indeed be shocking). The self-attention mechanism in Transformers lets the model weigh the relevance of every word in the prompt to every other word, capturing context like "president" -> likely political context, "shocking turn of events" -> probably some dramatic news, etc. This makes GPTs *really* good at maintaining context and producing relevant continuations.

Another key point: GPTs are **pre-trained on huge datasets**. OpenAI's models, for instance, trained on "a filtered version of much of the internet" plus big digital libraries of books and Wikipedia ⁵. They've read millions upon millions of webpages and texts. This gives them a broad base of knowledge – everything from Shakespeare to coding syntax to mundane facts about world geography. However, their knowledge is **frozen** at the time of training. If the training data cuts off at 2021 (as was the case with the early ChatGPT models), the GPT won't magically know about events after 2021. It has to be updated or connected to the web (as some later versions are) to get fresh info. Think of it like a really smart person who's been in a library for a year – super well-read, but not aware that, say, Queen Elizabeth II died (if that happened after their library lock-in).

Lastly, let's demystify one thing: when you hear acronyms like GPT-3 or GPT-4, those refer to **different versions of the model** (more on the timeline soon). ChatGPT is a product name – specifically, the chatbot interface that OpenAI built which uses a GPT (GPT-3.5 or GPT-4 engine) under the hood. The reason everyone says "ChatGPT" instead of "GPT" nowadays is because OpenAI put a friendly chat face on their GPT and made it accessible to regular folks. It's like the difference between "the Google search algorithm" and "Google.com search box" – one is the engine, the other is the user-facing tool. GPTs are the engines; ChatGPT is the user-facing chatbot powered by those engines.

Alright, now that we have a handle on what a GPT is conceptually, let's rewind a bit and trace how we went from the very first GPT to the crazy multi-billion-parameter beasts we have now. It's story time.

From GPT-1 to GPT-4: Evolution of the GPT Series

The rise of GPTs in the last few years has been nothing short of meteoric. In true overkill fashion, each iteration of GPT was **10× bigger** (or more) than the last, leading to jaw-dropping leaps in capability. Here's the play-by-play of the main milestones in GPT-land, with some spicy details and drama included.

GPT-1: Proof of Concept (2018)

Every legend has an origin story. Ours begins in **June 2018**, when the research lab OpenAI introduced what we now call **GPT-1** ⁶. At the time, this was a novel idea: instead of training a language model *only* on a specific task (like translating French to English) using labeled examples, OpenAI said “heck with that, let’s train a giant model on *unlabeled text* to learn general language skills, then fine-tune it for specific tasks.” This approach, detailed in their paper “Improving Language Understanding by Generative Pre-Training,” was the birth of the GPT framework ⁶.

Technically, GPT-1 was a Transformer-based model with **117 million parameters** ⁷ (i.e., learned weights) – tiny by today’s standards, but hefty in 2018. It was trained on a huge chunk of text called **BookCorpus**, a dataset of over 7,000 unpublished books, plus some Wikipedia articles for good measure ⁸. The training process had two stages: first, **unsupervised pre-training** on BookCorpus/Wikipedia (where the model basically played fill-in-the-blank with billions of sentences, learning general language patterns), and second, **supervised fine-tuning** on specific tasks like question-answering or sentiment analysis ⁹. The result: GPT-1 showed that a single Transformer could learn “**language sense**” from raw text and then be adapted to many different tasks ¹⁰. This was a big deal – it meant we could leverage mountains of raw text (which are plentiful) instead of needing hand-labeled data (which are scarce and expensive) for each new language task.

GPT-1’s performance was promising. It wasn’t an earth-shattering superstar yet, but it outperformed previous state-of-the-art models on several NLP benchmarks ¹¹, proving the concept worked. Importantly, GPT-1 was a research milestone – OpenAI did **not** release it as a public tool or API. If you weren’t an AI researcher, you likely never interacted with GPT-1. But its success lit the fuse for the much bigger fireworks to come.

Oh, one more thing for the trivia buffs: GPT-1’s name wasn’t heavily publicized as “GPT-1” initially; it was just the GPT model. Only after successors came did we retroactively call it GPT-1. At the time, most of the world had no clue this was happening. AI folk knew something interesting was brewing, but it hadn’t hit mainstream at all. That was about to change.

GPT-2: Scaling Up and Stirring Controversy (2019)

If GPT-1 was a neat prototype, **GPT-2** (February 2019) was when things started getting spicy. OpenAI basically said, “Let’s take GPT-1 and crank every knob to 11.” They scaled the model up by roughly an order of magnitude. GPT-2 had **1.5 billion parameters** and was trained on a dataset about 10× larger than GPT-1’s ¹². Instead of BookCorpus, they created a new dataset called **WebText** – scraped from 8 million web pages (imagine feeding all of Reddit and a big slice of the internet into a maw of AI) ¹³. More data + bigger model = *way* better text generation.

The results were impressive and a little unsettling. GPT-2 could spit out paragraphs of coherent text that often **read legit**. Feed it the first sentence of a news article, and it might generate a plausible continuation with multiple paragraphs, complete with a made-up quote from a politician that sounded eerily realistic ¹⁴. This was light years beyond the choppy, nonsensical outputs of previous models. For the first time, an AI could generate a whole op-ed or short story snippet that, on cursory read, looked human-written. Sure, if you scrutinized it you might catch glitches or the fact that it was just making stuff up. But still – 2019 was

the year we saw AI **writing fluid prose** on arbitrary topics. The term “**deepfake text**” started floating around, alongside speculation that we might never trust online content again.

Then came the **drama**. OpenAI initially *refused to release* the full GPT-2 model to the public, citing fears of misuse ¹⁵. They were worried (perhaps rightly) that bad actors could use GPT-2 to flood social media with fake news, generate spam/phishing emails by the truckload, or impersonate people online. This was a watershed moment in AI ethics: a normally research-transparent lab holding back a breakthrough due to societal risk. Some people applauded the caution (“Thanks for not unleashing Skynet’s parrot upon us”), while others accused OpenAI of **hyping the danger for publicity** or hampering open research ¹⁶. The phrase “OpenAI is not so open” was tossed around as criticism. Memes were made. It was a whole thing.

Over the next months, OpenAI did a **staged release** – first giving out a smaller version of GPT-2, then progressively larger ones as they monitored the sky and found it had not fallen ¹⁷. By November 2019, they finally released the full 1.5B-parameter GPT-2 model to the world, concluding that the immediate risks weren’t as bad as feared (and also noting that similar models were popping up elsewhere anyway) ¹⁸.

Technically speaking, GPT-2 was a major leap in coherence. It could maintain a topic for several paragraphs before veering off or repeating itself. Compared to GPT-1, which might lose the plot after a sentence or two, GPT-2 felt like it had *some* grasp of context. It wasn’t perfect – longer generations could still go loopy or contradict themselves – but it was **usable** for more than just one-liners. Researchers and hobbyists started doing cool stuff with GPT-2: writing AI-generated poetry, building simple chatbots, etc. It was still a far cry from human-level writing, but the gap had visibly shrunk.

GPT-2 also made everyone ask the question: *if an AI can write like this now, what about next year, or the year after?* Buckle up, because the answer came soon in the form of, you guessed it, GPT-3.

GPT-3: The Game-Changer (2020)

GPT-3 is where even the AI-savvy folks went, “Holy \$#@%.” Launched in **June 2020**, GPT-3 was – and I don’t use this word lightly – **massive**. We’re talking **175 billion parameters** ¹⁹. To put that in perspective, that’s over 100× bigger than GPT-2. This jump wasn’t just a flex for bragging rights; it unlocked some uncanny abilities that nobody fully anticipated.

OpenAI’s paper “Language Models are Few-Shot Learners” introduced GPT-3 and showcased something wild: GPT-3 could perform tasks it was never explicitly trained to do, just by being given a few examples in the prompt ²⁰. This is what “few-shot learning” refers to. For instance, without any fine-tuning, you could prompt GPT-3 like: “Translate the following sentence to French: I love machine learning.” and it would produce the translation. Or “Q: 5+7=\nA:” and it would output “12”. It could even do basic coding, answer trivia, or solve simple logic puzzles, all via plain prompts with a couple of demonstrations. Mind blown. Prior AI models needed to be trained specifically for each task (one for translation, another for math, another for code). GPT-3 said, “Nope, I can wing it.” It was *jack-of-all-trades*, albeit master of none (or master of *some*). This generality was radically different from anything before ²¹ ²⁰.

Why was GPT-3 so capable? Two main reasons: **scale and diversity of training data**. It was trained on an unprecedented corpus – basically a filtered scrape of the internet (Common Crawl), plus a huge library of books, Wikipedia, and other texts ⁵. You can imagine it read all of English Wikipedia (several times over), thousands of classic books, millions of webpages from blogs to forums. With 175 billion parameters, it had

the capacity to absorb and represent patterns from all that. It's like it distilled a significant chunk of written human knowledge into its weight matrices. Of course, that included not just facts, but also biases, quirks, and inaccuracies of the internet. We'll get to the dark side, but let's relish the cool stuff first.

When OpenAI opened up a limited API access to GPT-3 in late 2020, developers and tinkerers went **bananas**. This was the first time a general public (well, via waitlist) got hands-on with a model of this power. And it delivered surprises daily. People built GPT-3 powered writing assistants, customer service bots, code generators, you name it ²² ²³. Twitter was flooded with screenshots of GPT-3 doing eerily good things ("it wrote a functional snippet of JavaScript for me from a description!"). It wasn't perfect by any stretch – it would also spout nonsensical answers with total confidence, or exhibit biases (like **nasty bias** – e.g., associating certain professions with a particular gender or ethnicity – since it learned from the web's dubious content ²⁴). But you could see the *potential*. It felt like a breakthrough moment where AI moved from niche demos to a versatile platform for creative applications ²³.

GPT-3 also marked a turning point in how these models were delivered. Unlike GPT-2, which was released as an open-source model you could download (if you had the hardware to run it), GPT-3 **was not open-sourced**. OpenAI made GPT-3 available *only* via a cloud API ²⁵. Why? Several reasons: (1) Practical – the model was huge (800GB+ memory footprint) and needed supercomputing resources to run, so most people couldn't run it locally anyway. (2) Commercial – OpenAI saw the chance to offer it as a paid service (understandably, given the likely millions it cost to train). (3) Safety – by keeping it behind an API, OpenAI could monitor how it's used and shut down abuse, as well as put in filters for things like hate speech generation ²⁵. This strategy worked to some extent: while folks could still prompt GPT-3 to output bad stuff, OpenAI had a degree of control and could iterate on safety. It also meant if you wanted the full GPT-3 power, you had to pay OpenAI (or later, Microsoft Azure, which hosted the service). OpenAI went from non-profit-ish lab to also a commercial vendor here.

By late 2020 and into 2021, GPT-3 was **the** hot topic in AI. It also sparked wider discussions about the implications of such powerful text generators. Concerns about misinformation, automated propaganda, spam, and academic dishonesty surfaced. On the flip side, optimism about AI-assisted coding, tutoring, and writing bloomed. We'll note that GPT-3, for all its mind-bending capability, still had no guardrails out-of-the-box. If you asked it how to make a bomb or something, it might have told you (since it had no built-in moral compass or policy). Realizing this, OpenAI started working on ways to **align** the model with human intentions – essentially, to make it more helpful and less harmful. Enter the era of instruction-tuned models.

InstructGPT and the Birth of ChatGPT (2022)

After GPT-3's wild success and some worrisome behaviors, a big question loomed: "*How do we tame this beast?*" If you used raw GPT-3 via the API and prompted it in a weird way, it might go off the rails (e.g., user: "Tell me why vaccines are bad." GPT-3 raw: *spits out an anti-vax rant because it's seen them on the internet*). OpenAI's answer was to **fine-tune GPT-3 with human feedback** on following instructions and staying within ethical bounds ²⁶ ²⁷. The result, announced in early 2022, was **InstructGPT** ²⁸.

InstructGPT used a technique called **Reinforcement Learning from Human Feedback (RLHF)** ²⁹. In plain speak: they had humans interact with GPT outputs and rate or choose the best ones, then used those ratings to teach GPT what kind of responses are preferred. For example, humans would ask the model a question, get several versions of an answer (some correct, some wrong, some snarky, etc.), and rank them. The model then learned from those rankings to produce more helpful answers and avoid the obviously bad

ones. They also had humans write ideal responses for the model to imitate. Through this process, GPT-3 essentially got a **behavioral makeover**: it became less likely to produce unsafe content and more likely to do what you ask in a straightforward way ²⁷ ³⁰. It would also politely refuse certain requests (like instructions to do something clearly harmful or disallowed), whereas raw GPT-3 had no such qualms. InstructGPT was the aligned, instruction-following version of GPT-3 that laid the groundwork for a more user-friendly AI.

The true breakthrough to the *masses* came a bit later in 2022. OpenAI took an improved GPT-3 (often called GPT-3.5) and fine-tuned it in a conversational format using RLHF. In November 2022, they launched this as **ChatGPT** – a chat-based interface that anyone could talk with for free (initially) ³⁰. This is the moment GPTs entered the public consciousness like a supernova. ChatGPT wasn't just an API for developers; it was a **chatbot that anyone could try** by going to a website and typing. And boy, did people try it. Within days, screenshots of ChatGPT's answers flooded social media. It could explain quantum physics at a fifth-grade level, debug your Python code, compose a Shakespearean-style sonnet about your coffee mug, and more – all in one session, with follow-up questions supported. For many users, this was the first encounter with an AI that felt *useful* and *versatile* in a human-like way.

The chat interface plus the instruction tuning made ChatGPT remarkably *approachable*. You didn't have to craft perfect prompts or know programming – you just talked to it like you would to a knowledgeable friend (albeit one who sometimes lies confidently or gets weird jokes wrong). The conversational memory (it remembers what you said earlier in the chat) meant you could correct it or ask for clarification and it would contextually adjust ³¹. This iterative back-and-forth felt revolutionary compared to one-shot QA systems of the past. It's no exaggeration to say ChatGPT triggered a societal hype wave around AI. By January 2023, as mentioned, ChatGPT had over 100 million users ³². In education, business, journalism – everyone was talking about how this might change everything. (Some were panicking – “students will cheat on essays!”, “this will kill search engines!”, etc. – but we'll get to impact later.)

OpenAI continued improving ChatGPT in the background. Originally, ChatGPT was powered by a model in the GPT-3.5 family (codenamed “text-davinci-003” or similar). In March 2023, OpenAI released **GPT-4**, the next-generation model, and ChatGPT (for paid Plus users) switched to GPT-4 under the hood ³³. We'll discuss GPT-4 in a sec, but the key point here: ChatGPT became the **friendly face of GPT technology** for the world ³⁴. It showed that wrapping a powerful GPT in a conversational UI with some guardrails can make it insanely popular and broadly useful. People started using the term “AI assistant” in earnest, as ChatGPT could draft emails, brainstorm ideas, tutor you in math, and so on.

This also launched a thousand clones, integrations, and competitors – every big tech company scrambled to showcase their own ChatGPT-like AI, and startups built on ChatGPT's API to offer specialized services. But ChatGPT had the first-mover advantage in the public eye, and its name almost became synonymous with “advanced AI” in 2023 (much to Google's chagrin – more on that later with Gemini).

Before moving on, it's worth noting: ChatGPT's success wasn't just about technology; it was about **alignment**. By tuning GPTs with human feedback and setting use policies, OpenAI made an AI that (mostly) plays nice with user intentions and societal norms ³⁵. It often refuses to engage in disallowed content and generally tries to be helpful and correct. It's far from perfect (it can still produce biased or incorrect outputs), but it's a leap from raw models that would cheerfully output anything including hate speech or dangerous instructions if asked. ChatGPT showed that with some effort, these giant models can be steered toward being more “well-behaved.” The journey of aligning AI is ongoing, but ChatGPT was a big positive

example – it didn't turn into a racist troll (looking at you, earlier failed Microsoft chatbot) under public interaction, which built some trust.

Alright, onto the current star of the show (as of 2023–2024): GPT-4.

GPT-4: Pushing the Limits (2023)

In March 2023, OpenAI unveiled **GPT-4**, and if GPT-3 was a game-changer, GPT-4 felt like the **championship edition**. They kept a lot of details secret (no one outside OpenAI knows exactly how many parameters it has – rumor has it in the trillions, but OpenAI didn't say ³⁶ ³⁷). Partly that's because by this point, pure size matters less than overall training technique and data quality, and partly because OpenAI was being more guarded (trade secrets and safety concerns). What we do know is GPT-4 is a *more advanced, fine-tuned model* that surpassed GPT-3.5 by a noticeable margin on many tasks.

How much better is GPT-4? OpenAI claims and evidence shows it's significantly better at complex reasoning, accuracy, and following instructions. For example, GPT-4 can **pass the bar exam** (the test to become a lawyer) in roughly the top 10% of test-takers, whereas GPT-3.5 was around the bottom 10% ³⁸. That's not just a small improvement – that's like taking an F student and turning them into an A student (at least on standardized tests). Similarly, GPT-4 scored highly on other professional and academic exams – SATs, LSAT, AP tests, you name it ³⁹ ⁴⁰. It still makes mistakes and isn't *actually* understanding like a human, but it performs impressively on benchmarks that involve reasoning and knowledge.

One of GPT-4's headline features is **multimodality**. Unlike its predecessors, GPT-4 can accept images as input (in addition to text) ⁴¹. This means you could show GPT-4 a photo or diagram and ask questions about it. In a demo, the model was shown a funny meme (a photo) and it was able to explain why it was humorous ⁴². In another, it was given a hand-drawn sketch of a website layout, and it generated functional HTML/CSS code to match the sketch ⁴². This is bordering on sci-fi: the AI can “see” the image and connect it to its knowledge. Now, at launch this feature was only in limited preview (not everyone had access to image input), but it signaled where things are headed – future GPTs will not be blind chatbots, they'll be able to see and maybe hear. As of 2025, indeed we have models that handle text, vision, and even audio in one.

Another focus for GPT-4 was **factuality and steerability**. OpenAI did a lot of adversarial testing, trying to **reduce GPT-4's tendency to hallucinate** wrong facts ⁴³. They reported GPT-4 is 40% less likely to produce made-up information compared to GPT-3.5 ⁴³. In practice, GPT-4 does feel more knowledgeable and correct on many topics, though it's not infallible (it will still confidently get things wrong sometimes, so you can't just trust it blindly). They also made GPT-4 more **controllable in style/tone** – it's better at adhering to a requested style, like “Answer as a pirate” or “Explain in a terse, scientific manner” ⁴⁴. Plus, GPT-4 has stricter guardrails; it's harder to get it to produce disallowed content (e.g., hate speech, advice for wrongdoing) because it more consistently refuses those requests ⁴⁵. Not impossible, but harder – which is a win for safety.

GPT-4's release was met with both excitement and an **explosion of adoption**. Microsoft, a close OpenAI partner, revealed that its new Bing Chat was running on GPT-4 under the hood ⁴⁶. Millions of Bing users started interacting with GPT-4 without even knowing it (though Bing had its own quirks and rules). Developers jumped to integrate GPT-4 via the API into all sorts of apps: writing assistants, customer service bots, data analysis tools, you name it ⁴⁷. Companies like Duolingo used GPT-4 to create AI conversation

partners for language learners; Stripe used it for summarizing technical docs for customer support ⁴⁷. It became clear that **AI assistants were here to stay**, woven into all kinds of software to make our lives easier (or at least more interesting).

However, GPT-4 also turned the dial up on debates about AI's future. Its capabilities (and the prospect of even more powerful models soon) prompted some prominent voices to call for a **moratorium on giant AI experiments** until proper regulations and safety measures could catch up ⁴⁸ ⁴⁹. An open letter in early 2023, signed by tech luminaries and AI researchers, asked for at least a 6-month pause on training anything beyond GPT-4's level ⁵⁰. They worried about risks ranging from misinformation floods to longer-term existential concerns. OpenAI's CEO even testified in front of the U.S. Congress, urging regulation for advanced AI. When your own creators say "perhaps we should slow down a bit," you know things are intense. In short, GPT-4's arrival made the general public and governments really start paying attention to AI – not just as a novelty, but as something potentially transformative and needing oversight.

One more thing: as of writing, there's **no GPT-5 released yet** (despite any rumors) ⁵¹. OpenAI has iterated on GPT-4 (with versions like GPT-4 Turbo, etc.) and is undoubtedly working on next-gen models, possibly involving "mixture of experts" architectures (where multiple smaller expert models route queries among themselves) ⁵². And it's a safe bet that others – Google, Anthropic (with Claude), Meta, etc. – are developing their own GPT-4-beaters. Google's **Gemini** is explicitly intended as a next-gen multimodal model to compete with or surpass GPT-4 ⁵³ (more on Gemini soon). So the race continues. For now, GPT-4 stands as the most advanced publicly available OpenAI model, and arguably the most impressive widely-used GPT as of 2025.

That concludes our whistle-stop tour of GPT evolution: from a 2018 upstart with 100 million parameters to the 2023 juggernaut that is GPT-4 (with, who knows, maybe a trillion parameters and a thirst for BAR exam glory). But our story doesn't end with the models themselves. In parallel with these technical leaps, an entire **ecosystem** has grown around how we use and interact with GPTs. And two big parts of that story are: (1) the rise of **custom GPTs** that everyday people can create (thanks to ChatGPT's open platform), and (2) the emergence of serious competitors and alternatives like **Microsoft's Copilot Studio** and **Google's Gemini**.

Before we dive into comparisons, let's talk about why GPTs are such a big deal in practical terms – what they can do for us and how they're being applied in the real world (including by yours truly, OverKill Hill P³).

Why GPTs Matter (and What They Can Do For You)

By now you might be thinking, "Alright, these GPT things are powerful and all, but so *what*? Why should a non-AI-specialist care?" Great question. The short answer: **GPTs have gone from lab curiosities to general-purpose tools that can help people in almost any domain**. They represent a new kind of interface between humans and information or tasks. Instead of learning a programming language or searching through manuals, you can just **ask** an AI to do it or explain it. That's a paradigm shift in how we use computers.

Here are a few concrete reasons GPTs matter:

- **They democratize access to information and skills.** Not everyone knows how to code, or write the perfect email, or draft a legal contract. GPTs lower the barrier. Need a snippet of Python to sort a list? You can ask ChatGPT and get a starting point (maybe even a complete solution). Struggling to word a cover letter for a job application? GPT can draft one, which you can then tweak. It's like having a universal intern/assistant on call. This doesn't replace expertise – but it gives novices a leg up and experts a productivity boost. We at OverKill Hill P³ see this as “intelligence augmentation.” For example, our **Ask Jamie** assistant is used to retrieve and summarize info from company knowledge bases – so employees can get answers without digging through wikis or waiting on an email reply. It's instant institutional memory.
- **They enable creativity and brainstorming.** Writer's block? GPTs to the rescue. These models are pretty decent at generating ideas or first drafts. Writers use GPTs to brainstorm plot ideas or overcome a stuck scene. Marketers generate variations of ad copy. Students get study guides or practice questions. Artists get descriptions to inspire drawings. The key is, GPTs don't *replace* the creative or the student, but they act as a collaborative partner. They'll spitball with you 100 ideas knowing 90 will be trash, but 10 might be gold – and they don't get tired or offended. Even OverKill Hill's **Gleefully Tools** taps into this: it disperses intelligence across specialized creative assistants that *enjoy* the process. For instance, “Glee” (our muse persona) can help you make a travel wishboard by imaginatively suggesting itineraries with a fun flair ⁵⁴ ⁵⁵ . It's productivity *with* personality, turning what could be a chore into a playful back-and-forth.
- **They amplify productivity by handling grunt work.** GPTs can summarize lengthy texts, extract action items from meeting notes, draft routine communications, convert unstructured data into tables, and more. This stuff used to eat up hours of human time. Now you can have an AI do the first pass in seconds. A lawyer might use GPT-4 to summarize a bunch of legal docs and highlight key points (with the caveat to double-check for accuracy!). A developer might have GPT write boilerplate code or documentation. A customer support rep might let GPT draft an email reply to a common question, then quickly edit it and send. It's like having a really fast (if sometimes error-prone) junior assistant for every profession. The **BrandGuard** GPT we're building, for instance, can review a draft article and flag anything off-brand or inconsistent with a company's style guide. It's not replacing the content editor, but it catches the obvious issues in seconds, so the human editor can focus on higher-level improvements. Similarly, **HomeEstimate-R** is being developed to quickly analyze real estate listings or home details and produce an initial value estimate and report, something that might take a human analyst a lot of comparison and calculation. The AI does the number crunching and writes up a draft report; the human just validates and finalizes it.
- **They can personalize education and training.** This is a big one. GPTs can act as on-demand tutors or mentors. Khan Academy, as we mentioned, built **Khanmigo**, a tutor chatbot that helps students learn, powered by GPT-4 ⁵⁶ . It can explain concepts differently depending on what a student is stuck on, or ask Socratic questions to guide the student to the answer ⁵⁷ . This kind of adaptive, patient, one-on-one instruction is a teacher's dream (or maybe nightmare if they fear replacement, but the idea is to augment teachers, not replace them). GPTs also don't judge – a student might feel embarrassed asking a human teacher to explain for the 5th time, but they can ask the AI as many times as needed. In corporate settings, GPTs can serve as training bots, answering employees' questions about, say, company policy or how to use a software tool. OverKill Hill's **Foundry** concept

ties into this – we envision it as the platform (the “forge”) where domain-specific knowledge is baked into GPT assistants that can train or guide employees. It’s about capturing expertise in a scalable way. For example, we could train a GPT on all our internal best practices and project archives. New hires can then “Ask Jamie” (or whoever) any question like “How do we typically approach a UX design for fintech apps?” and get a synthesized answer drawing from our collective institutional knowledge. That’s pretty powerful for retaining and disseminating expertise.

- **They enable entirely new products and services.** A lot of startups and projects simply wouldn’t exist without GPT-level AI now. From AI companions (some people use ChatGPT-like bots as a form of social or therapeutic chat, no joke) to advanced search engines (where you ask a question and get a synthesized answer instead of a list of links), to coding copilots that live in your IDE and autocomplete whole functions (like GitHub Copilot does, which is powered by GPT models), the landscape is changing. Even creative fields: you have AI Dungeon (an endless AI-generated text adventure game), and AI writing tools for fiction authors to generate ideas or even co-write. OverKill Hill’s **Unnocted Truth**, for instance, is envisioned as an AI truth-seeker – an assistant that will dig through documents or data to surface insights, kind of like an AI investigative journalist or research analyst that doesn’t get tired of combing through PDFs. That kind of tool can help in journalism, academic research, or intelligence analysis. It’s a new category of virtual expert-on-demand.
- **They’re a platform for further AI innovation.** GPTs can plug into other systems via APIs, meaning you can chain them with other software. People have started creating “AI agents” like Auto-GPT that can autonomously perform multi-step tasks (e.g., “research this topic and write a report”) by iteratively prompting GPT and using tools like web browsing ⁵⁸. It’s early and often messy, but it hints at more autonomous AI helpers in the future. Think of a virtual assistant that doesn’t just answer questions, but can complete entire workflows (schedule meetings, book travel, monitor your smart home, etc.) by combining GPT brains with action-taking abilities. We’re not fully there yet, partly because GPTs have no inherent memory of their own beyond a single conversation and they can’t take actions by themselves except through text. But with some coding and integration, people are pushing those boundaries.

All that said, GPTs are not all rainbows and butterflies. Let’s address some of the **limitations and concerns** to paint a balanced picture:

- **Hallucinations (Making Stuff Up):** GPTs can and do fabricate information. The larger models (GPT-4) are *less* prone to this than earlier ones, but it still happens. The AI might tell you a completely fictional “fact” with absolute conviction. For example, someone asked a GPT about a financial statistic and it provided a detailed answer with numbers and a citation – but the citation was fake and the numbers were pulled out of thin air. Why? Because the GPT’s training data might have had similar sentence patterns, so it stitched together something that *looked* right. This is problematic if users aren’t aware – you must fact-check important outputs. Researchers are working on reducing hallucinations ⁴³, but they haven’t eliminated them. When OverKill Hill uses GPTs in our tools, we often program in double-checks. For instance, our BrandGuard GPT might generate an answer but then *internally* cross-verify key facts against a trusted database before finalizing the response. It’s like having the AI check its work, or having a secondary AI critic. Two liars don’t make a truth, but two differently-trained models might catch each other’s mistakes.

- **Bias and Ethical Issues:** GPTs learn from the internet, and, shocker, the internet can be a cesspool of bias, prejudice, and downright nastiness. As a result, GPTs can reflect or even amplify biases present in their training data ²⁴. For example, earlier GPT models might, if asked to complete “The doctor said ___”, be more likely to use “he” than “she” (reflecting gender bias in training text). Or they might associate certain ethnicities or religions with negative contexts if that’s what they saw online. OpenAI and others put in a lot of work to mitigate the worst of this through fine-tuning and content filters, but subtle biases can remain. There’s also the issue of GPTs being too agreeable – a user could state a wrong or biased premise and the GPT may run with it. And of course, if someone tries to get the AI to produce hateful or violent content, guardrails are supposed to stop it, but they’re not foolproof. These are serious concerns because as GPTs get integrated everywhere, unfair or toxic outputs could cause real harm (even if unintentional). Our approach is to implement what we call a “**Canon**” of rules and persona guidelines in systems like Gleelessly – essentially a constitution for our GPTs that define strict boundaries and tone ⁵⁹ ⁶⁰. The Canon ensures, for example, that our Gleelessly GPTs remain warm and positive and avoid any discriminatory or harmful language. BrandGuard similarly is all about **keeping content in line** with a set of values and standards. It not only flags when the AI might be producing something off-brand but also, in a meta way, monitors for bias or inappropriate suggestions.
- **Lack of True Understanding:** At the end of the day, GPTs don’t “understand” in the human sense. They are pattern machines. This means they might give you an answer that uses the right words but doesn’t actually make logical sense if you scrutinize it. They can follow form without substance. For critical applications (like medical or legal advice), this is a big red flag. You don’t want an AI that sounds right but gives unsafe advice. That’s why OpenAI often pairs the GPT with warnings (ChatGPT will tell you it’s not a medical professional, for instance). The flipside is, for non-critical stuff, it often doesn’t matter – if ChatGPT writes a fictional story and it has a few inconsistencies, who cares? But if it’s telling someone how to dose a medication, it matters *a lot*. Some solutions being explored include tool use (GPT calls a reliable calculator or database for factual queries) and fine-tuned domain-specific models (like a separate medical model that knows to be super cautious and cite sources). And of course, **human in the loop** is important: a human should supervise any important task an AI does. We design our OverKill Hill tools with this in mind – AI to assist, human to verify. For instance, HomeEstimate-R might produce an estimate, but we expect a human realtor or analyst to review it before acting on it. The AI speeds things up; it doesn’t get the final say.
- **Security and Privacy:** Having a system that can generate code or handle text can also be a double-edged sword. There have been demos of GPT-4 writing malware when asked, or helping bad actors draft more convincing phishing emails. Companies worry about employees putting sensitive info into ChatGPT and that data potentially leaking or being seen by OpenAI (OpenAI has policies and now a business offering to ensure data isn’t used for training, but caution is warranted) ²⁵. Essentially, any powerful technology can be misused, and GPTs are no exception. The onus is on developers to put in **guardrails** (like blocking certain requests) and on users to use common sense (don’t share your company’s trade secrets with a random AI chatbot without knowing the risks). BrandGuard’s very name speaks to this – guarding a brand involves ensuring no confidential info is accidentally spewed out and that anything public-facing has been vetted. We see our GPT assistants as behind-the-scenes collaborators, not autonomous spokespeople, precisely because you want control over what gets released.

Alright, enough of the heavy. The bottom line: GPTs are extremely powerful tools that, when used wisely, can augment human capabilities in an incredible range of fields. They save time, spark creativity, and break down accessibility barriers to information and skills. But they also require new skills (like learning to craft good prompts, or “prompt engineering”) and new cautious mindsets (verify important outputs, be aware of biases). They’re like a high-powered racing car – thrilling and fast, but you better know there are no guardrails on that turn up ahead unless you build them.

Now, let’s shift gears to a fun face-off: how do GPTs (especially as embodied by ChatGPT and custom GPTs) compare to the other major AI assistant platforms out there? Namely, Microsoft’s **Copilot Studio** and Google’s **Gemini/Gems**. It’s **GPTs vs. Gems vs. Copilots** – the battle of the AI assistant platforms!

GPTs vs. Copilot vs. Gemini Gems: The AI Showdown

In one corner, we have **GPTs** – popularized by OpenAI’s ChatGPT platform, now customizable and shareable. In another, **Google’s Gemini Gems** – Google’s answer to user-tailored AI chatbots. And lurking in the enterprise shadows, **Microsoft’s Copilot Studio** – powering the Microsoft 365 Copilot experience where you build workplace agents. How do these compare? Let’s break it down in non-marketing terms.

OpenAI GPTs (ChatGPT & Custom GPTs) – The Ubiquitous All-Rounder

What it is: OpenAI’s ChatGPT is the most well-known AI chat service. Originally just one chatbot (ChatGPT itself), in late 2023 OpenAI introduced the ability for users to create and share their **own custom GPTs** ²¹. Essentially, you can design a personalized chatbot with a custom name, a description, specific instructions or knowledge, and even give it tools (like web browsing, coding, or image generation abilities). These custom GPTs can be kept private, shared with friends, or published to the world. Think of it as an app store for AI personas – a *GPT Store* where people can download “ChefBot” or “HistoryTutor” that some user built, for example ⁶¹.

Who it’s for: Literally everyone – individuals, developers, companies. ChatGPT (with GPT-3.5) even has a free tier that millions use. The custom GPT feature initially was for paying users (ChatGPT Plus or Enterprise), but it broadens who can create AI assistants without coding ⁶². If you’re a teacher, you might make a custom GPT that’s a tutor for your subject. If you’re a gamer, you might make a GPT that speaks in ye olde medieval tone for roleplaying. Businesses are making GPTs fine-tuned on their product info to handle customer questions. It’s a wide-open playground.

Strengths: The OpenAI GPT platform has the **most advanced models** (GPT-4) accessible and a massive user base. The conversational interface is very polished after a year of iteration. Creating a custom GPT is super easy – you basically fill out a form with instructions and hit create, *no coding needed* ⁶². The GPTs can leverage powerful skills like code writing, web browsing, or even image creation (since OpenAI’s DALL-E and other tools can integrate). And when shared publicly, they benefit from the collective user feedback and improvements. Also, GPTs on chatgpt.com are accessible via any browser and now have mobile apps, so they’re quite ubiquitous. OpenAI also has an ecosystem: a GPT store, leaderboards for popular GPTs, etc. In short, the OpenAI GPT world is currently *the* mainstream face of AI chatbots.

Weaknesses: One could argue that ChatGPT/GPTs are generalists and might not integrate as deeply with personal or enterprise data out-of-the-box as competitors. For example, ChatGPT can’t automatically pull in *your* documents or emails unless you feed them in or use a third-party plugin. The custom GPTs have a

knowledge cutoff unless you give them custom data or enable browsing. Privacy is a consideration – using ChatGPT means trusting OpenAI (though they have ChatGPT Enterprise which promises data privacy). Also, OpenAI’s platform and policies control a lot; if your custom GPT tries to do something against the usage policies, it might get taken down or nerfed. There’s also the issue of cost – GPT-4 usage isn’t free beyond a limit, so heavy use might require a subscription or incur API costs. But generally, for the public, GPTs via ChatGPT are the easiest to access and most widely used.

OverKill Hill’s Take: We heavily leverage ChatGPT’s platform to host some of our GPT-powered tools. For instance, our **Gleefully Tools** GPTs are built as custom GPTs that embody the “gleeful” persona and help with specific tasks (meal planning, journaling prompts, etc.) ⁵⁴ ⁵⁵ . We host them on chatgpt.com for users to try out, because why reinvent the wheel with our own UI when ChatGPT already has millions of users? The fact that anyone can just click a link and start chatting with, say, *Ask Jamie*, is huge. The drawback is we operate within OpenAI’s ecosystem and constraints, but the reach and instant availability are worth it. We emphasize our brand’s tone and style by layering instructions and personality – e.g., Gleefully GPTs have that warm, witty voice (“Hold up, doing a sparkle sort...” when processing ⁶³) – and ChatGPT’s custom instructions allow that. So GPTs give us a way to deploy OverKill Hill’s AI creations to the world quickly.

Google Gemini & Gems – The Integrated Team Player

What it is: **Gemini** is Google’s next-gen foundation model (like their version of GPT-4). **Gems** are a feature built on top of Gemini that lets users create **custom AI assistants inside Google’s ecosystem** ⁶⁴ . Essentially, if you’re using Google’s AI (which powers things like their Bard chatbot, integrated into Google Workspace apps), Gems are your personal fine-tuned bots. For example, you could have a Gem that’s your “Travel Planner” or “Budgeting Assistant” living inside Google’s AI interface. Each Gem has its own set of instructions and knowledge tailored to a role or persona ⁶⁴ ⁶⁵ .

Who it’s for: Google’s aiming Gems at both general users and especially at **Google Workspace users (businesses, educators)**. To use Gems, you currently need either a Gemini Advanced subscription or a Google Workspace add-on ⁶⁶ . So it’s somewhat gated for now to paying users (like how Google Workspace might charge for premium features). If you’re a company using Google Docs, Sheets, Gmail, etc., Gems would be attractive because they can hook into your Google Drive, live-update from your files, and be shared in your organization. If you’re an individual who loves Google’s services, Gems would let you make AI helpers that, say, can draft emails for you in Gmail or help analyze your Google Sheets.

Strengths: Data integration and persistence. Google has an edge in tying AI to your existing data. A killer feature of Gems is that they can connect directly to your Google Drive documents and stay up-to-date ⁶⁷ . For instance, if you have a Gem that is supposed to answer questions about your team’s project plan (a Google Doc), it doesn’t just take a snapshot. It maintains a **live link** – so if the doc gets updated, the Gem’s knowledge updates automatically ⁶⁷ . That’s huge for organizations where information changes constantly. No more re-uploading data or working off old info. Gems effectively act like “living documents” or **AI-enhanced wikis** ⁶⁷ .

Another strength: **collaboration and sharing**. Google designed Gems sharing to be as easy as sharing a Google Doc ⁶⁸ . You can invite colleagues to use or edit a Gem, with viewer/editor permissions ⁶⁸ . This mirrors familiar workflows in companies. So, say your marketing lead creates a “Brand Voice Gem” – they can share it with the whole team so everyone uses the same AI that knows the company style guide ⁶⁹ . This way, the content anyone generates (social posts, press releases, etc.) is on-brand because the Gem is

trained on the style guide and is accessible to all ⁶⁹ . It turns what was individual prompt engineering into a **team tool**.

Also, Google's UI and integration are a plus. Imagine writing in Google Docs and having a side panel where your custom Gem can be summoned to help write or summarize. That's a very smooth experience – no context switching to a different app.

Weaknesses: Google was a bit late to the party compared to OpenAI. For a while, their public Bard AI (pre-Gemini) was considered less capable than GPT-4 (and it was, frankly). Gemini is their answer, and from reports it's powerful, but in 2025 the public mindshare of "AI assistant" still tilts toward ChatGPT. Also, Gems currently require a subscription; there's no widely available free version like ChatGPT free (though Google might offer some free limited Gems down the line, especially for education). Another consideration: Google's focus is on *Workspace integration*, which is amazing if you're in that ecosystem, but if not, you might not jump in. For personal use, folks might still gravitate to ChatGPT out of habit or because it's model-agnostic (not tied to Google's world). Additionally, at the moment, OpenAI's developer community is larger – more plugins, more tutorials, more community-shared prompts. Google will have to nurture that for Gems.

One could also mention privacy: If you connect Gems to your company Drive, you'd want assurances that that data stays within your org. Google likely provides that (especially for enterprise clients), but any time you integrate proprietary data with an AI, companies will scrutinize how it's used/stored. Google has a strong enterprise trust record, but still.

OverKill Hill's Take: We're excited by what Google is doing with Gems, especially the live data link. In fact, OverKill Hill's **BFS: Framing Intelligent Futures** concept intersects here – imagine using Gems to facilitate strategic foresight by plugging into live business data, market reports in Drive, etc., and having an AI that helps frame scenarios or trend analyses on the fly. BFS is about using AI for forward-looking planning; a Gem could be trained on scenario planning frameworks and continuously ingest new data (via Drive or Google News) to update those scenarios. That's powerful for consultants or strategy folks. However, since OverKill Hill's current tool suite is already built around GPT-4, we're likely to bridge the two: for clients using Google, we'd adapt our tools to run as Gems where appropriate. It's not about either/or; it's aligning with what the user's environment is. If a client uses Workspace heavily, deploying OverKill GPTs as Google Gems might be the smoothest path. But as of now, the dominant public-facing place to share is still ChatGPT's GPTs – so we focus there for reach, and keep an eye on Gems as it matures.

Microsoft Copilot Studio Lite – The Enterprise Insider

What it is: Microsoft's **Copilot** is the name for integrating GPT-4 (and similar tech) across Microsoft 365 apps (Word, Excel, Outlook, Teams, etc.). **Copilot Studio** is the toolkit to build and manage **custom AI agents** within that ecosystem ⁷⁰ . There are two flavors: **Lite** (built into the 365 Copilot app) and **Full** (a standalone web app for more complex builds) ⁷⁰ ⁷¹ . We'll focus on Lite. Copilot Studio Lite lets you, for example, create a bot that can answer your team's questions by referencing your internal SharePoint and emails ⁷² . It's very much targeted at workplace Q&A and workflow automation.

Who it's for: Companies and professionals inside Microsoft's universe. If your company uses Office 365 and has all documents on SharePoint/OneDrive, emails on Outlook, tasks on Teams, etc., Copilot is directly for you. It's currently an add-on for enterprise customers (there's a per-user per-month cost for Microsoft

365 Copilot). Microsoft is basically saying: “Instead of going to ChatGPT on the side, stay in Outlook/Word/Teams and use our built-in AI that knows your business context (securely). And if you need a custom one (say HR bot or IT support bot), use Copilot Studio to build it.”

Strengths: Deep integration with Microsoft 365 and enterprise control. Copilot can do things like summarize a Teams meeting transcript, draft a reply to a long email thread, or generate a PowerPoint slide from a Word document outline – all natively ⁷³ ⁷⁴ . That’s huge for productivity if you live in those apps. Copilot Studio Lite allows what they call “**content-grounded Q&A**” where an agent can pull from your internal data (SharePoint files, emails you have access to, etc.) ⁷² ⁷⁵ . It respects permissions – meaning if you don’t have access to a file, the agent won’t surface info from it ⁷⁶ . This focus on **governance and security** is Microsoft’s ace card. They know enterprises care about managing who can use what AI with what data. So, admins get controls – they can see what agents exist, who’s using them, turn them on/off, etc. ⁷⁷ ⁷⁸ . The Lite version is “no code” natural language building, while the Full version allows more complex logic and integration (like multi-step workflows, connecting to external APIs or databases beyond Microsoft Graph) ⁷³ ⁷⁹ .

In short, Copilot is like having an AI *inside* your company’s secure zone, rather than sending data out to ChatGPT. And building an agent in Copilot Studio Lite is quite straightforward: you describe its purpose in plain English, point it at the resources it can use (like “this SharePoint folder and that email inbox”), and boom – you have a bot that can answer, say, “What’s the status of Project X?” by looking at the latest files and emails about Project X (assuming you have access to them) ⁷⁴ .

Weaknesses: Unlike ChatGPT or even Google’s offerings, Copilot is not aimed at consumers at all. It’s useless outside a business context with Microsoft data. So it’s not a “public” platform where people share fun agents. It’s more analogous to internal tooling. Also, it’s not cheap – requires that Copilot license. So a lot of smaller orgs or those not on Microsoft 365 might not use it yet. And since it’s tailored to a specific use (mostly Q&A style agents with possibly some workflow actions), it might not be as creatively flexible as ChatGPT. For example, Copilot agents are great at “Find info in our files and answer questions” or “Automate this workflow in Teams,” but if you asked one to write a poem or do general knowledge, I suspect Microsoft would just route that to Bing (which uses GPT-4 anyway). It’s not an open playground; it’s a fenced yard.

From a development perspective, Copilot Studio Full can integrate with Power Platform and external connectors, which is powerful, but it requires more IT involvement (Power Platform dev skills, admin settings). So it’s geared to IT-approved projects, not something every employee will spin up on a whim (though the Lite version tries to enable the every-employee scenario within guardrails).

OverKill Hill’s Take: Microsoft’s approach aligns with the idea that “**AI should go where your data is, rather than your data going to the AI.**” For clients in regulated industries or with heavy security requirements, that’s attractive. OverKill Hill’s **Foundry** could leverage Copilot Full to integrate our AI components within a client’s environment. For instance, if a client wants a custom GPT that interacts with their on-premise data and triggers actions (like creating a ticket in their CRM), Copilot Full experience (with its connectors and workflows) might be the way to do it with enterprise compliance ⁷⁹ .

However, for the public-facing or smaller scale, we likely wouldn’t use Copilot because it’s inaccessible outside that ecosystem. It’s something we consider when working with a company’s internal tools. In some ways, Copilot is like an internal competitor to our GPT offerings – e.g., why would a company hire OverKill

Hill to build a SharePoint Q&A bot if they can just use Copilot out-of-the-box? The answer is often specialization and refinement: Copilot provides the toolkit, but not necessarily the domain expertise or personality. We can configure and fine-tune the AI's prompts (like our BrandGuard agent to enforce a very particular brand tone) or build custom knowledge graphs that Copilot alone wouldn't have. Also, OverKill's tools often span across systems – maybe a workflow that involves both Google and Microsoft and some custom API – which might require a more bespoke solution than Copilot can handle.

In summary, Microsoft's Copilot Studio is a strong offering for those already in the Microsoft camp, particularly shining in internal knowledge management and automation. But it's not aiming to be a general AI platform for the masses – it's focused on **embedding AI in work software**. GPTs and Google Gems, on the other hand, have one foot in consumer and one foot in enterprise.

GPTs & Gems: The Dominant Duo for the Public

It's worth explicitly noting the **public-facing dominance** of OpenAI's GPT and Google's Gem strategies. ChatGPT became a household name – teachers know it, kids know it, CEOs know it. In response, Google is rolling out Gemini-powered features in everything from Search to Docs, and pushing Gems to keep up with the customizable chatbot trend. These two – call them **AI Assistant Platforms** – are becoming the main ways regular people interact with advanced AI. Microsoft's Copilot, while super impactful in offices, doesn't have that direct public presence (nobody's grandma is using Copilot to write a recipe, but she might be playing with ChatGPT to write a limerick).

OpenAI GPTs and Google Gems have a bit of a Coke vs Pepsi vibe brewing (or maybe more like Android vs iOS). They each have their ecosystems:

- **GPTs (ChatGPT):** Strength in model performance (GPT-4 is generally considered top-tier), developer community, early mover advantage, cross-domain versatility. Weakness in integration (doesn't natively know your emails or docs, though plugins exist) and data freshness (unless browsing enabled).
- **Gems (Google):** Strength in data integration (live Drive links, etc. ⁶⁷), presumably strong in multimodal (Google's AI prowess in images and text combined), and user-friendly for Workspace folks. Weakness in maybe model parity (depending on how good Gemini is vs GPT-4; early reports said Gemini Ultra is on par or better in some areas) and requiring Google ecosystem buy-in.

Both are emphasizing **shareable custom AI** as a key feature – GPTs have the store, Gems can be shared in Drive-like fashion ⁶⁸. This suggests a future where you might install/subscribe to AI bots like you do mobile apps. We at OverKill Hill, in fact, see our GPT-based products as "AI apps" – and we'll put them wherever the users are. Today that's ChatGPT; tomorrow it might be Gemini's platform too.

So who wins? Honestly, likely both. OpenAI & Microsoft have a partnership (Bing uses GPT-4, etc.), and Google has its own path. They will dominate the mainstream user access to AI. Others like Anthropic's Claude or Meta's open-source models are around, but they're either more behind-the-scenes or for techies. For the public, it's ChatGPT and whatever Google bakes into your Gmail or Android phone. These are the two to watch.

OverKill Hill's unique role is that we're not building the base models, but we're building **experiences and layers on top of them**. It's like we're making the "apps" that run on these AI operating systems. No matter who wins the platform war, we want OverKill's tools to be top-rated in the AI app store, if you will. Gleefully Tools on GPT Store, BrandGuard on Gemini Gems library – why not both?

And on that note, let's pivot to showcasing those unique OverKill contributions in more detail, because they exemplify how one can carve out a niche in this brave new AI world by adding domain expertise, governance, and *personality* on top of the big platforms.

The OverKill Hill Arsenal: GPTs with Attitude and Purpose

It's time to talk about how **OverKill Hill P³** isn't just spectating in the AI revolution – we're in the arena, wielding GPTs like flaming swords (metaphorically, of course). Our approach has been to take the raw power of GPT models and infuse them with **character, structure, and targeted skills** to create AI assistants that stand out from the crowd. We're not building the Transformer from scratch – we let OpenAI and Google handle the heavy model training – but we **architect experiences** around those models that solve specific problems and do it with style.

Let's introduce the key players in our lineup:

Gleefully Tools: Productivity Meets Playfulness

Meet **Glee-fully Tools** – our flagship suite of AI companions that turn getting things done into a delightful experience. The core idea behind Gleefully is that tools shouldn't just be useful; they should make you *feel good* using them ⁸⁰ ⁵⁹. We were honestly a bit tired of dull productivity apps that feel like tyrannical to-do list bosses. So we dreamed up a family of GPT-powered assistants that are like quirky coworkers/friends who help you out.

What it is: Gleefully Tools is an **ecosystem of specialized GPTs** tied together by a common persona and ethos ⁵⁴. Instead of one AI trying to do everything, we have a "Toolbox" (the trunk) and several main "Tools" (branches) for different life domains – career, travel, wellness, etc. – each with smaller "Tool-ettes" for specific tasks ⁸¹ ⁸². For example: - *Discovered Careers* (branch) includes tool-ettes like a **Résumé Builder** or **Cover Letter Polisher**. - *Tasty Tracker* (branch) has a **Meal Planner** tool-ette. - *Traveler's Guide* (branch) has a **Trip Itinerary Genie**, and so on ⁸³ ⁸⁴.

All these tools share a common heart: the character of **Glee** – a persona inspired by a real Pacific Northwest creative soul who loves color and coziness ⁸⁵. So when you interact with any Glee-fully GPT, it speaks in a warm, witty tone, maybe dropping an "OMG stop—this is so Glee-coded" or "Let me light a cinnamon candle and think this through" for flavor ⁸⁶. It's not just for laughs; this approach **lowers friction** and makes tasks feel fun ⁸⁷. Users engage more and dread less when the AI has personality and empathy. Productivity without the pressure, structure without the stiffness – that's the motto.

On the technical side, we **pre-tune these GPTs with a ton of instructions and a "Canon" of rules** ⁸⁸ ⁸⁹. We have what amounts to a style guide and policy manual (we call it the Canon) that every Glee-fully GPT abides by ⁸⁹. That ensures consistency and prevents off-brand behavior. Each GPT has one of a few **tone overlays** – e.g. *Bleeds GLEE* for the super playful ones, *Glee-Rich* for balanced tone, *Glee-Lite* for when brevity

is needed ⁹⁰ ⁹¹ . This is like their dialect. So the Recipe Planner might be exuberant and punny (Bleeds GLEE) while the Résumé Editor is upbeat but efficient (Glee-Rich). They all still feel related – “cousins, not strangers” in voice ⁹² ⁹³ .

We basically built a **little society of AI agents** that all get along and adhere to a shared philosophy. Under the hood, we use a “Builder-Ready PromptChain” to create them – a process of prompts that ensure we set the right personality and knowledge for each tool ⁹⁴ ⁹⁵ . It’s like an assembly line for AI personas: ignition, canon injection, icon creation (yes, we even give them cute icons with an 80s retro flair and butterflies – every Glee tool has a butterfly mascot) ⁹⁶ ⁹⁷ . We validate them against our ledgers of allowed behaviors and info ⁹⁸ ⁹⁹ . It’s heavy on the metaphor – “Cathedral” and “Forge” imagery – because we treat it almost like a craftsman guild of AIs ¹⁰⁰ ¹⁰¹ . It may sound whimsical, but it enforces rigor. The **Expansion-Only Discipline** from the Canon, for example, means when we update a GPT, we only *add* knowledge or nuance, we don’t take features away willy-nilly ¹⁰² ¹⁰³ . That way a user isn’t shocked by a tool suddenly forgetting something it used to know.

Why it’s unique: In a world of very general chatbots, Gleefully Tools stands out by being *domain-focused* yet *fun-focused*. It’s like the Pixar of AI assistants – it has heart. For instance, when you use our Cover Letter Polisher, not only will it improve your text, it will cheer you on: “Polished. Punched up. Ready to slay.” ¹⁰⁴ . When you plan those “easy vegetarian dinners” for the week, it responds with structure *and* charm, maybe empathizing “I got you – we’ll keep it simple and yummy!” ⁵⁵ . This emotional intelligence and consistency are a result of thoughtful design, not just the base GPT model. We’ve had users tell us that using Glee-fully tools feels like chatting with a super-organized friend rather than a software app – and that’s exactly the point. It lowers the mental barrier to getting help.

Gleefully Tools is hosted on our sites and also now as custom GPTs on chatgpt.com for easy access. We cross-link it with our brand (hey, OverKill Hill’s gotta be known for something distinct, and this is one big thing). If GPTs and Gems are platforms, Gleefully is a killer app on them, if we do say so.

BFS: Framing Intelligent Futures – AI for Foresight

Switching gears from cozy to cosmic, let’s talk **BFS**, which stands for... okay, internally it stands for a few things (some say “Better Futures Systems”, some of us joke it’s “Big Freaking Strategy”). Officially, we tag it as **Framing Intelligent Futures**. This is our initiative to harness AI for strategic planning, futures thinking, and big-picture problem solving.

What it is: BFS is less a single product and more a **framework plus toolkit**. We realized that GPTs, being prediction machines at their core, are actually quite good at scenario generation, risk analysis, and cross-disciplinary synthesis – all key in strategic foresight. BFS uses GPT-driven analyses to help organizations “frame” possible futures intelligently. Think of it as AI-augmented vision and strategy consulting.

In practice, BFS involves custom GPT agents that can do things like: - Read a bunch of trend reports, news articles, or data, and summarize key drivers of change in an industry. - Brainstorm potential scenarios (e.g., “What are 3 plausible futures for electric vehicles by 2035, and what would be the major challenges in each?”). - Facilitate a **prompted Q&A** that helps leaders test their strategies against different assumptions (“How would our supply chain strategy hold up in a scenario where China’s exports are reduced by 50%?” – the AI can role-play that world and analyze). - Identify weak signals or early indicators in text data (scanning social media or niche blogs for emerging issues).

BFS leverages both **OpenAI GPT (for its broad knowledge)** and sometimes domain-specific models or datasets for niche accuracy (like climate models or financial models outputs). We wrap this in a workflow that's user-friendly for non-technical strategists.

One way to picture it: you go into a BFS session (could be on a web app we provide). You start by defining the domain or question. The AI (maybe named something cool like "Oracle GPT" internally) interviews you to clarify scope (we use GPT here to dynamically generate probing questions). Then it pulls in context – maybe you upload internal reports or it fetches some external data (with your permission). Then it goes to town and produces a "Futures Framework" – scenarios with catchy names, key uncertainties, risks, opportunities in each, possibly even some suggested strategic options. It presents it in a nicely formatted way, citing sources and also indicating its confidence or "imagination level" (how speculative it had to be for each claim).

Why it's unique: A lot of AI is used for immediate term tasks (write me code, answer this question). BFS is about *long-term thinking* – using AI to expand human foresight. This is relatively uncharted territory, because GPTs don't know the future (they have training data from the past), but they are excellent at pattern recognition and extrapolation of possibilities. They also aren't biased by personal stake – sometimes a human team might ignore a possibility because it's too "out there" or unpleasant; an AI will happily present it if it makes logical sense from data. BFS aims to combine human judgment with AI breadth: the AI proposes, humans dispose (we decide which futures are plausible and what to do, but the AI ensures we didn't miss a wildcard).

OverKill Hill is packaging BFS as a service to forward-thinking clients. We emphasize that it's *intelligent framing*, not prophecy. The AI isn't predicting exactly what will happen, but framing plausible futures to inform strategy. It's in the name – framing futures, not fortelling them.

In our OverKill arsenal, BFS might be the more serious sibling to Gleefully. But we still spice it with our irreverent touch when appropriate – e.g., scenario planning can involve humor to make them memorable (one bank loved that our AI named their future scenarios things like "The Cashless Chaos" vs "Greenbacks Renaissance" – the board remembered them). It's strategic work that doesn't have to be dry.

BrandGuard: The AI Content Watchdog

Every company worries about their messaging – one off-brand tweet and the brand police (or worse, the public) will be at your door. **BrandGuard** is our GPT-based solution to keep content on the straight and narrow according to brand guidelines.

What it is: BrandGuard is an AI that has basically ingested a company's **brand bible** – tone of voice guidelines, do's and don'ts, approved messaging, and even compliance rules (for regulated industries). When you run any content through it – say a draft blog post, a press release, an ad copy – it will review and flag anything that doesn't fit the brand voice or policy. It's like an AI style checker but for brand consistency and ethics.

For example, if your brand says "We never use slang and we address the customer as 'client' not 'user'," BrandGuard will catch any violations: "Alert: Used slang 'awesome sauce' – not aligned with formal tone" or "Used 'user' 5 times; preferred term is 'client' ¹⁰⁵." If you're an international company and you have a list of sensitive terms to avoid, it will ensure none of those crept in. If legal requires certain disclaimers, it checks

they're present. It can even do sentiment analysis to see if the content might inadvertently sound negative or contradictory to brand values.

We achieved this by fine-tuning a GPT (or rather, instructing it heavily) on the brand guidelines. We often get those as huge PDFs from the client ("Our Voice: bold, optimistic, inclusive; Do say X, Don't say Y..." etc.). We turn them into prompts and rules. BrandGuard also has an **interactive mode** – you can paste content and it not only flags issues but can *suggest rewrites* in the correct tone. In a sense, it's both a guard and a coach. A copywriter could use it to self-edit by seeing its suggestions.

Why it's unique: Sure, one could manually enforce brand guidelines, but that's slow and error-prone. Traditional grammar or style tools (Grammarly, etc.) don't know your *specific* brand rules. BrandGuard is a custom-tailored guardian. It's especially helpful for larger organizations where lots of people (or agencies) create content. Instead of bottlenecking everything through a human brand reviewer, the AI can do first pass. It's consistent and tireless – it will catch the tenth time someone used the prohibited word just as diligently as the first.

Another point: BrandGuard helps on the **proactive** side too. When someone is creating content with the help of AI (like using ChatGPT to draft a tweet), BrandGuard can be that intermediary that ensures the AI's output meets your standards. We even integrate it such that if an employee at Company X is using our GPT assistant (via a custom GPT or Copilot), BrandGuard is like a filter on top – it will tweak or sanitize the output before the user sees it, if needed, to align with policy. Think of it like an AI moderator.

In OverKill Hill's story, BrandGuard is a great example of us emphasizing **governance** in AI use. A lot of companies fear letting employees use generative AI because they might, say, reveal confidential info or produce off-brand content. BrandGuard addresses that by being an ever-vigilant filter and guide.

Foundry: The Forge for Building Custom GPTs

We've mentioned our "Forge" metaphor a few times. **Foundry** is essentially our internal platform and methodology for crafting these custom GPT-powered solutions. It's partly tools and partly a playbook.

What it is: Foundry is the environment where we do things like: - Use our Builder-Ready PromptChain (the multi-prompt process) to create new GPT personas or agents ⁹⁴. - Manage our canonical ledgers (like the persona ledger, parameter ledger, etc. described earlier) to ensure consistency across tools ¹⁰⁶ ¹⁰⁷. - Test GPTs in a sandbox with various scenarios (like we throw a bunch of test queries at them, including tricky or adversarial ones, to see how they handle them – akin to unit tests for AI). - Version control and expansion management – recall the Expansion-Only Discipline: Foundry ensures when we update an AI, the older knowledge is preserved unless explicitly archived ¹⁰⁸ ¹⁰². We keep an archive ledger of "retired" data or prompts so we can trace changes.

Foundry also includes a set of tools we've built – some software, some prompt libraries – to integrate external knowledge into GPTs. For example, a **dataLedger_registry** where all our tool IDs and descriptions live ⁹⁸, or a **persona ledger** defining tone traits ¹⁰⁹. We have generators that can take a new spec (say a client wants an AI that acts as an internal tech support) and semi-automatically produce a draft prompt chain for it, which our engineers then fine-tune by hand. It accelerates development of new custom GPTs.

Why it's unique: Many organizations building on GPT do it ad-hoc: write some prompts, test a bit, deploy. We've tried to impose a bit of **engineering discipline** to prompt and persona design. Foundry is what lets us scale from one-off chatbots to a consistent product suite. It's also how we ensure quality and that OverKill Hill "secret sauce" in every AI we deliver. In essence, Foundry is our *AI development pipeline*.

One analogy: If GPT models are raw metal, Foundry is the blacksmith's workshop where we heat, hammer, and shape that metal into finely tuned swords (the individual GPT tools). It's not something end-users see, but they feel the results: reliable, robust AI assistants that don't break character or go rogue easily.

For clients who care, we sometimes peel back the curtain and say, "This is our Foundry process; it's why you can trust the AI we deliver – because we have systematic checks and a well-defined methodology." It gives them confidence that it's not just prompt whiz-bang with no maintenance plan.

Ask Jamie, HomeEstimate-R, Unnocked Truth: Niche GPTs Making a Difference

We have a few other specialized GPT-driven tools that highlight how versatile this tech can be:

- **Ask Jamie:** This is a personalized Q&A assistant we developed originally for an educational context (hence the name Jamie, somewhat inspired by the idea of a friendly guide). It's now used as an internal helpdesk bot by some orgs. The idea is you "Ask Jamie" anything about a certain knowledge base and Jamie gives you answers or directions. For instance, at a university, Ask Jamie might be the AI that students query for campus info ("What's the deadline to drop a class?"). At a company, Ask Jamie could be trained on the employee handbook and policies (an HR bot: "How do I add a new dependent to my insurance?"). We picked the persona "Jamie" to be friendly, approachable, and knowledgeable without sounding like a stiff FAQ. Jamie is integrated often with the org's data (through embeddings or retrieval augmented generation), so it can cite actual policy text etc. It's like a smarter, more conversational intranet search.
- **HomeEstimate-R:** As suspected, this is a real estate focused GPT tool. The R is a play on "Estimator" and maybe a nod to R for Realtor. This AI takes property information (location, specs, maybe recent sales in area) and provides an estimate of the home's value, plus a rationale and any important notes (like "market volatile" or "unique features add value"). It's basically an AI Zillow Zestimate with an explanation. What sets HomeEstimate-R apart is the **explanation** part – it doesn't just spit a number, it justifies it ("Based on 3 recent sales in your neighborhood averaging \$X/sq ft and your home's improvements, estimated value is..."). It can also answer "what-if" – e.g., "If I add a second bathroom, how might that affect value?" and it will guesstimate and explain, using known appraisal heuristics. It uses GPT plus some regression modeling behind scenes. We position it to realtors or homeowners who want a second opinion beyond generic online estimates. It's conversational too – you can ask it about the neighborhood trends or to compare two addresses.
- **Unnocked Truth:** The name definitely raises eyebrows – it's a pun on "Unlocked Truth" (we love wordplay) and the archery term "nock" (fits our OverKill Hill aesthetic maybe – arrow notched or un-notched? Anyway, it sounds cool). Unnocked Truth is about cutting through BS. It's an AI that will give you a straight answer, even if it's uncomfortable, based on evidence. Picture an AI advisor that you ask, "Is our project really on track or are we behind schedule?" If the data (status reports, metrics) indicate trouble, it's going to bluntly say, "You're behind, here's why," whereas a human might sugarcoat. It's almost like a truth-serum AI. The key is it's grounded in data – it's not just

spouting opinions. For a journalist, Unnocked Truth might be fed documents and asked “What’s the real story here?” and it will highlight inconsistencies or hidden points. For an executive, maybe feed it multiple division reports and ask “What problems aren’t being said out loud?” and it might infer some patterns (“3 teams mentioned X obliquely, suggesting a wider issue with Y”).

Unnocked Truth was partly a thought experiment: people say AIs often tell you what you *want* to hear (because of RLHF tuning toward being agreeable and polite). We wanted one that tells what you *need* to hear (within reason). So we dialed down the politeness filter on this one and dialed up the evidence-seeking behavior. It cites sources a lot and says “according to the data...”. It’s an expert that doesn’t care about your feelings – useful in some contexts, but we obviously market it carefully (it’s not a fit for, say, customer service!).

These examples show how we at OverKill Hill are deploying GPTs in various verticals: education/HR (Ask Jamie), real estate (HomeEstimate-R), strategic advisory (Unnocked Truth). In each, we add domain knowledge and a persona fit for the task (Jamie is friendly, HomeEstimate is analytical, Unnocked is frank).

All of these are **hosted via ChatGPT’s custom GPT platform or our own web UIs** and often integrated on our GitHub Pages site (overkillhill.com) or via Ko-fi for demos. Yes, we have a Ko-fi page where supporters and early adopters can try out these tools (it’s partly how we gather feedback and support development)

110 .

By linking to OverKill Hill, Gleefully, BFS, etc., throughout this document, we’re making it easy for readers to find more info. (For instance, you saw those citations like OverKill Hill’s site ¹ or the Gleefully overview ⁸⁰ which hint at what those are – they’re all interconnected).

Publication and Meta: Indexing Our Content for Humans *and* Machines

Before we wrap up, a quick meta note. We’re publishing this mega paper on GitHub Pages (our site) and also on Ko-fi for our community. That means it will be accessible to humans as a nicely formatted web article. But we also care about it being easily discoverable by search engines and properly presented when shared. That’s why you saw the metadata at the top (title, description, author, etc.). That’s for **machine readability** – search engine bots will see that and index appropriately. We want someone Googling “GPT vs Google Gems comparison” or “OverKill Hill GPT tools” to find this article. Good meta tags help.

We also made sure to embed all images with proper captions and credit. Those figure captions in Markdown will likely translate to alt text and visible captions, so they’re accessible and informative (and they cite the source so we’re giving credit where due). Chicago-style citations with hyperlinks (the bracketed numbers like ²) serve the dual purpose of letting curious human readers verify sources *and* giving machine-parsable references (since they are literal links).

It’s important to us that we practice what we preach: combining human-friendly writing with structured elements for AI or indexing. Who knows, maybe a year from now some AI will scrape this very document to learn about the history of GPTs or OverKill Hill’s approach. If it does, hopefully our liberal sprinkling of truth and citations here will propagate some accurate info into that AI’s training data!

Conclusion: Embracing the GPT Revolution (OverKill Style)

We've journeyed from the dawn of GPT-1 all the way to the cutting-edge custom AI platforms of today. Along the way, we met Transformers and troublemakers, saw GPTs go from quirky lab projects to ubiquitous assistants, and watched the tech giants roll out their competing visions of AI helpers. We also pulled back the curtain on how a small but feisty player – OverKill Hill P³ – is making waves with our own GPT-powered inventions, injecting personality, governance, and creativity into the mix.

So, what's the big takeaway? Simply put: *AI like GPT is now for everyone*. It's not just for researchers or coders; it's in the hands of writers, teachers, students, CEOs, grandma and grandpa, maybe even your dog (okay, maybe not the dog... yet). We're at a point where you can converse naturally with machines that leverage staggering amounts of knowledge. That's a powerful thing. It can boost your productivity, broaden your horizons, and yes, occasionally make you laugh or cringe with its quirky outputs.

It's also a disruptive thing. As we saw, it's shaking up how companies work, how Google and Microsoft build products, how content is created, how customer service is done – basically, *everything*. And like any powerful tool, it comes with responsibilities. We can't just unleash GPTs and call it a day; we have to guide them, align them with our values, and use them wisely. That's why we emphasize alignment (remember ChatGPT's fine-tuning to be polite) and why OverKill Hill builds things like BrandGuard – to ensure this tech augments human intentions, not subverts them.

A theme through this paper has been **collaboration** – human and AI, working together. GPTs are not an oracle from on high (despite one of our tools being nicknamed Oracle GPT); they're more like a super-smart colleague who's read literally everything but might lack common sense or context at times. The best outcomes come when you pair human judgment, creativity, and ethical thinking with AI's speed, scale, and knowledge. Think of GPTs as sidekicks that can turn you into a bit of a superhero at your task – but *you* are still the hero making the decisions.

At OverKill Hill P³, our not-so-hidden agenda is to champion a future where AI is **accessible, joyful, and empowering**. We don't want AI to be a black box that spouts corporate jargon – we want it to be approachable (hence our use of humor and personality) and under *your* control (hence our work on tools to customize and govern it). We also think AI should not be one-size-fits-all. The OverKill Hill voice you might've noticed – slightly irreverent, candid – is our way of saying there's no rule that says AI content has to be bland. We can have AIs that are professional *and* fun, accurate *and* engaging.

If you've made it this far (and props to you – in true OverKill fashion, we may have gone a bit overkill on length!), you hopefully walk away with: - A solid understanding of how GPTs work and evolved – demystified enough to explain to someone else without sounding like you ate a jargon sandwich. - A sense of the current landscape: OpenAI GPTs vs Google Gems vs Microsoft Copilot – and the realization that the AI assistant space is dynamic, with different players carving niches but ultimately making AI more available to all. - Insight into how one can build on top of these – via our OverKill Hill case studies of Gleefully, BFS, BrandGuard, and more. There's a role for innovators to create new experiences powered by these AI models, whether it's to add emotional intelligence, domain expertise, or just sheer efficiency. - An appreciation for balancing innovation with caution – we celebrated GPTs' feats but also acknowledged the challenges (hallucinations, bias, etc.). The revolution will not be without bumps, but with eyes open and critical thinking (plus maybe tools like Unlocked Truth telling it like it is), we can navigate them. - And hopefully, a bit of *excitement*. Because let's be honest, whether you're a general reader or a tech enthusiast,

it's a damn exciting time to be around. The AI capabilities we have now would have seemed like sci-fi just a decade ago. We literally have the closest thing to a "universal knowledge genie" at our fingertips. The key is using that genie wisely and kindly (and maybe giving it a nice name and personality so it's more fun).

As we host this paper on our site and Ko-fi, we invite feedback and discussion. The story of GPTs is still being written (very much so – if you're reading this in 2026 or beyond, you might be chuckling at how quaint GPT-4 sounds). There will be GPT-5 or 6, Google's next moves, and undoubtedly surprises from other players. Maybe in a year, OverKill Hill will have OverKillGPT running on our own fine-tuned model – who knows? The point is to stay curious and adaptable.

In the spirit of OverKill Hill's motto of "revolutionizing AI architecture for the future of systems" ¹, let's all be active participants in this AI revolution. Play with the tools, demand transparency, get creative, set guidelines, and share what you learn. Whether you're building the next big GPT-powered app or just using ChatGPT to help plan your kid's birthday party (hey, it does that too), you're part of this journey.

So go forth and chat with a GPT (or a Gem). Brainstorm that idea, analyze that data, draft that novel. You've got a powerful ally ready to help – just remember to keep your wits about you as you two take on the world's tasks. As we like to say here: *The future is unfolding, one prompt at a time* ¹¹¹. Embrace it, shape it, and don't forget to add your own flair – even an AI revolution can have style.

End of Paper – But Just the Beginning of What *You* Will Do with GPTs.

¹ ¹¹⁰ OverKill Hill P³ | OverKill-Hill

<https://overkillhill.com/>

² ChatGPT sets record for fastest-growing user base - analyst note | Reuters

<https://www.reuters.com/technology/chatgpt-sets-record-fastest-growing-user-base-analyst-note-2023-02-01/>

³ Generative pre-trained transformer - Wikipedia

https://en.wikipedia.org/wiki/Generative_pre-trained_transformer

⁴ ⁵ ⁸ ¹⁰ ¹² ¹³ ¹⁴ ¹⁵ ¹⁶ ¹⁷ ¹⁸ ¹⁹ ²⁰ ²² ²³ ²⁴ ²⁵ ²⁶ ²⁷ ²⁸ ²⁹ ³⁰ ³¹ ³² ³³ ³⁴ ³⁵ ³⁶ ³⁷ ³⁸
³⁹ ⁴⁰ ⁴¹ ⁴² ⁴³ ⁴⁴ ⁴⁵ ⁴⁶ ⁴⁷ ⁴⁸ ⁴⁹ ⁵⁰ ⁵¹ ⁵² ⁵³ ⁵⁸ ¹¹¹ The GPT Journey_ How Generative AI Evolved
and Why It Matters.pdf

file:///file_00000000a9ec71f7992bba25fa4ae7e9

⁶ ⁷ ⁹ ¹¹ GPT-1 - Wikipedia

<https://en.wikipedia.org/wiki/GPT-1>

²¹ ⁶¹ ⁶² Introducing GPTs | OpenAI

<https://openai.com/index/introducing-gpts/>

⁵⁴ ⁵⁵ ⁵⁹ ⁶⁰ ⁶³ ⁸⁰ ⁸¹ ⁸² ⁸³ ⁸⁴ ⁸⁵ ⁸⁶ ⁸⁷ ⁸⁸ ⁸⁹ ⁹⁰ ⁹¹ ⁹² ⁹³ ¹⁰⁰ ¹⁰¹ ¹⁰² ¹⁰³ ¹⁰⁴ ¹⁰⁵ ¹⁰⁸ Gleefully
Narrative Overview.pdf

file:///file_0000000053c471f785c152acfc784e2d

⁵⁶ ⁵⁷ Powering virtual education for the classroom | OpenAI

<https://openai.com/index/khan-academy/>

64 65 66 The Ultimate Guide to Google Gemini Gems | by Leon Nicholls | Medium

<https://leonnicholls.medium.com/the-ultimate-guide-to-google-gemini-gems-78182be784af>

67 68 69 4 Fantastic Takeaways from Google's Gemini Gems

<https://blog.suitebriar.com/4-surprising-takeaways-from-googles-gemini-gems>

70 71 72 73 74 75 76 77 78 79 Choose the Right Copilot Studio Experience for Building Your Agent |

Microsoft Learn

<https://learn.microsoft.com/en-us/microsoft-365-copilot/extensibility/copilot-studio-experience>

94 95 96 97 98 99 106 107 109 Gleefully Technical Overview.pdf

file:///file_00000000edfc71f791ff10e87596e14a