



Coordinating Multiple GPTs: Current Capabilities and Limitations

Using Multiple GPTs Together (the “Mesh” Concept)

OpenAI has begun rolling out a beta feature that lets you use **multiple GPTs in one conversation** via Slack-style mentions ¹ ². In practice, this means you (the user) can **@mention** different custom GPT bots within a single chat, effectively switching which GPT is responding at each turn. Each custom GPT might be specialized (with its own instructions or knowledge base), so this setup allows a form of “swarm intelligence” where you divide tasks among expert agents ² ³. For example, you could have one GPT focused on research and another on writing, and call on each when appropriate ³.

However, keep in mind the reality of today’s system: **all these GPTs are still backed by the same underlying model (e.g. GPT-4)** – the difference is in their instructions and attached data. They don’t truly run in parallel or talk to each other autonomously; you as the user orchestrate the interaction by choosing which GPT to invoke. In essence, it’s like having multiple personas of GPT-4 available and you direct questions to one or the other as needed. This can certainly enhance productivity (e.g. sending a sub-question to a specialized GPT within the same chat ⁴), but it’s not a magic hive-mind – it’s a **manual “divide and conquer” strategy** enabled by the new mention feature ³. Each response still comes one at a time, from whichever GPT you tagged. There’s no simultaneous multi-agent dialog unless you explicitly prompt them in sequence.

Sharing Knowledge and Files Across GPTs

Yes – you can give multiple custom GPTs access to the **same reference documents or logic** by uploading duplicate copies of files to each GPT’s knowledge base. ChatGPT’s Custom GPT Builder allows up to **20 knowledge files per bot** ⁵, so you could attach the same set of files (containing your external logic, commands, or domain knowledge) to all the GPTs in your “array.” This means each GPT has the *same background information* to draw from when answering. In practice, the model *integrates the content of those files as background knowledge* rather than quoting them verbatim ⁵. It will refer to the files **before responding** to you, using them like an internal knowledge base via a retrieval-augmented approach ⁶.

Critically, these files are **read-only context** – the GPT won’t modify them, it will only use them to inform its answers. Each custom GPT “remembers” its attached files for every query, so you don’t have to keep re-uploading or reminding it of that content during the conversation. By giving all GPTs the same command definitions or reference data, you ensure they share a common logic and facts. For example, if you have a file of standard operating procedures or a schema, every GPT with that file will adhere to those instructions. This addresses your idea of shared external logic: in essence, the GPTs have *identical copies* of the rules/knowledge to follow. Just note that any time you update those files, you’ll need to **detach and reattach** the new versions for each custom GPT (there’s no live sync between them, as each has its own stored copy).

Practical tip: Keep your knowledge files **well-structured and succinct**. The model will use them more effectively if the information is clear and not overly verbose or redundant ⁷. Large or disorganized files can slow down responses or lead to missed info ⁷. In other words, you're doing the right thing by revising the schema files as needed – just make sure to update all GPTs with the new file when you do.

Persistent “Tether” to Knowledge Files

A custom GPT's attached knowledge files serve as a persistent context throughout the thread. In reality, the model doesn't “forget” these files as the conversation continues – they're part of the GPT's configuration and are consulted on each response ⁶. This is unlike normal ChatGPT, which can forget details between chats or once the context window is exceeded. Here, as long as you remain in a conversation with that custom GPT, it will always have its knowledge base available to draw from (no matter how many turns the conversation takes). In effect, there is a **strong tether to those files** for the duration of the chat. The system ensures the model *refers to the knowledge base before answering* your query ⁶, so important facts or definitions from your files should consistently inform its responses.

That said, keep expectations realistic: the GPT will use the files **contextually**, meaning it will pull in relevant info when it's applicable rather than dumping the whole file content. If a discussion veers far off-topic from anything in the knowledge base, the model might not mention the files at all (because nothing in them was relevant to that topic), but the tether is still there in the background. Moreover, the model is still bounded by a context window (for GPT-4, typically on the order of thousands of tokens). Very long conversations can push some earlier dialogue out of scope. The knowledge files, however, aren't just treated as one-time context at the start – they remain consistently accessible. In summary, you can trust that **the thread won't lose the attached knowledge** on its own. The files act like a permanent reference book the GPT consults every time, rather than something it might forget after a few replies.

Handing Off a Conversation Between GPTs (@mentions and Handover)

With the multi-GPT mention feature, you can indeed start a thread with one GPT and then **handover** to another by tagging it. Because it's all within a single chat session, the conversation history is shared context. In practical terms, when you switch to GPT-B (by typing `@GPT-B` and your question), GPT-B will see the prior messages in the thread (including what GPT-A said earlier). Each GPT will, of course, apply its own instructions/knowledge to the situation, but they **inherit the conversation context** seamlessly. Since you plan to give both GPTs the same knowledge files and command definitions, they also share the same background info – this helps make the transition smooth. GPT-B won't be confused about any domain-specific terms or logic that GPT-A was using, because GPT-B has the same reference material available.

In current usage, users report that this kind of hand-off *does* work. For example, one person described having several specialized internal GPTs and using the new mention ability to send a sub-question to a different GPT “in one chat” – essentially exactly what you're aiming to do ⁴. From your perspective, it will feel like you have a panel of experts and you're directing questions to the one best suited for each task, all while keeping a unified conversation thread. As long as you explicitly invoke the right GPT for the right question, the **context carries over** and the responses should follow logically from where the last one left off.

Do keep in mind a few realities of such handovers: the different GPTs may have different styles or roles, which means when you switch, the tone or approach of the answer might change (since each GPT follows its own instructions/persona). That's normal, and often the very reason to use multiple GPTs (each has a unique function or perspective). Also, GPT-B will only know what's been said or provided *in the shared conversation* plus its own knowledge files – it doesn't literally have GPT-A's internal state beyond what was communicated. For example, if GPT-A did some reasoning or had an assumed plan not written out, GPT-B won't magically know that unless it was mentioned. In practice, this just means you might need to ensure any critical info or decisions from GPT-A's turn are either in the conversation text (so GPT-B sees it) or you remind GPT-B of them when switching. Usually, though, since GPT-A's answer is visible in the chat history, GPT-B can pick up from that. Overall, the hand-off can be “**clean**” as you put it – you are effectively just changing the respondent, while the story so far remains in view.

Returning to the Original GPT (Mentions and Thread Originator)

Yes – you can always **@mention the original GPT** (or any other) again in the same thread. The term “spawning GPT” suggests the one that initiated the conversation. In the new system, that initial GPT is just one of the bots available to mention. So after you've switched to GPT-B for a while, you could come back with a prompt like `@GPT-A: ...` to bring GPT-A back into the dialogue. There's nothing preventing you from **alternating between GPTs** as needed. Think of it like addressing different people in a group discussion by name – you direct a question or statement to one particular bot by tagging it, but you're never “losing” the others. They're all present in the conversation and can be called on when their input is needed.

One thing to clarify is that **GPTs themselves won't autonomously call or mention each other** (at least not currently). Only the user can direct the flow to a different agent by using the @ syntax. So, you as the moderator will decide when to hand the mic back to GPT-A or over to GPT-C, etc. The interface will show you which GPT is responding each time (since you invoked it), but otherwise the chat just continues in one scroll. To summarize: you can certainly loop back to the GPT that started the thread or involve any others you've created, all within that one conversation. This flexibility is exactly what the mention feature is designed for – to let a team of GPTs collaborate under your guidance ⁸.

Reality Check: Constraints and Feasibility Today

It's great that you're pushing the frontier of GPT usage, but let's ground things in current **possibility and plausibility**. The multi-GPT setup you envision – essentially a mesh of GPT agents sharing knowledge – is **possible in a basic form today**, thanks to ChatGPT's custom GPTs and the new mention capability. You can set up multiple GPTs with identical knowledge files and coordinate them in one thread to tackle complex tasks in parts. This is very much *bleeding-edge* (a newly introduced feature), so expect some quirks and the need for manual oversight. Here are some key realities to keep in mind:

- **User-Orchestrated Interaction:** At present, GPTs won't spontaneously coordinate without your prompts. You have to direct which GPT should handle which query. The vision of fully autonomous agent swarms is on the horizon, but today the system relies on you to delegate tasks to each bot at the right time ² ³. Think of yourself as the project manager assigning questions to different expert assistants.

- **Same Core Model:** All your custom GPTs on ChatGPT Plus are built on GPT-4 (assuming you use the GPT-4 version). Their “intelligence” comes from the same place, just focused by different instructions or data. So while they can emulate a team of specialists, they are not truly independent minds with separate training – which is why one can’t surprise another with something outside of what GPT-4 knows or can do. They also share the model’s general limitations (for example, difficulty with extremely complex math or certain reasoning pitfalls, unless guided).
- **Context Limitations:** The conversation (plus any retrieved knowledge) is limited by GPT-4’s context window (typically ~8K tokens for standard GPT-4, unless you have the extended 32K version). This means if your thread becomes very long, the model might not explicitly recall early messages *word-for-word*. The knowledge files mitigate this somewhat by allowing the model to fetch relevant info on the fly ⁶, but even they won’t be fully utilized if the content isn’t triggered by your query. In practical terms, for lengthy sessions you may need to summarize or remind the model of key points occasionally. The good news is that the **attached files stay available** no matter how long the session, so core reference data remains at hand even if some earlier user messages scroll out of view.
- **Accuracy and Reliance on Knowledge Base:** By providing a solid knowledge base, you greatly reduce hallucinations, but it’s not a 100% cure. The model will try to use the files, but if asked something ambiguous or outside their scope, it might still “guess.” Make sure your prompts explicitly cue the GPT to use the files for factual questions (e.g., “According to our policy document, what is the procedure for X?”). The system is designed to refer to the knowledge base *before responding* ⁶, which is a huge advantage, just remember it uses the info to formulate an answer rather than quoting it unless you ask. Always verify critical outputs against your source files.
- **Performance Considerations:** Loading many large files can slow down the response. You mentioned a schema file – that’s fine, but if you have a lot of heavy documents attached, you might notice some latency. One user noted that having 13 very long documents made their custom GPT noticeably slower ⁷. Keeping files **concise and relevant** will help. It sounds like you plan to curate those files carefully (updating schema versions, etc.), which is the right approach for optimal performance and accuracy.
- **Deploying to Market:** Using ChatGPT Plus’s native features is fantastic for prototyping and even internal use, but deploying a consumer-facing product with this setup might be tricky. There is currently no way to embed the ChatGPT UI with multi-GPT functionality into your own app for external users (aside from directing them to use your custom GPT through the ChatGPT interface). If “deploy to market” is your end goal, you may eventually need to use the OpenAI API to recreate this multi-agent orchestration in a more controlled environment. The API doesn’t yet have a one-shot multi-GPT orchestration feature, so developers typically implement such logic themselves (having one GPT call another, etc.). The **good news** is that what you’re doing now – refining prompts, knowledge bases, and the interplay of specialized GPTs – will directly inform how you design an automated solution later. Just be aware of the platform’s terms of service and capabilities if you plan to commercialize; ChatGPT’s web interface is meant for individual use, whereas the API is intended for integrated products.

In conclusion, you *are* operating at the cutting edge of what ChatGPT can do right now. The idea of a GPT mesh or a fleet of AI assistants working in tandem is no longer pure hyperbole – it’s emerging reality, but

with **important constraints**. Today, you **can** array multiple GPTs with shared knowledge to simulate a collaborative team, but you'll need to drive that collaboration consciously. The system's "physics" (context limits, single-model backend, etc.) mean you won't get mystical synergy beyond what the model and your provided data allow. Still, this approach is plausible **today** and is likely to improve over time. Keep experimenting within these boundaries, and you'll be well-prepared for more advanced multi-agent features as they arrive. Good luck with building, testing, and eventually deploying your GPT-powered solution – it's an exciting space to be in, as long as you temper ambition with a clear view of the current limits!

Sources:

- OpenAI community discussion of the new multi-GPT mention feature ² ⁴
- AI news on GPT "@mention" capability rolling out to ChatGPT (multi-agent orchestration) ¹ ⁹
- Reddit Q&A on using knowledge files in custom GPTs (integration as contextual background, up to 20 files) ⁵ ⁷
- BYU GenAI guide on Custom GPTs (knowledge base is consulted before responding, akin to RAG) ⁶

¹ ⁸ ⁹ GPT Tags are Here. Enabling Using Multiple GPTs within a GPT Chat | by Saygin Celen | AI Frontier X | Medium

<https://medium.com/ai-frontier-x/gpt-tags-are-here-enabling-using-multiple-gpts-within-a-gpt-chat-aaa7cda99a11>

² ³ ⁴ New GPT beta feature allows you to chat with multiple GPTs in one chat. : r/singularity

https://www.reddit.com/r/singularity/comments/1acmmuq/new_gpt_beta_feature_allows_you_to_chat_with/

⁵ ⁷ Usage of knowledge files when creating a customGPT using the gptBuilder : r/ChatGPTPro

https://www.reddit.com/r/ChatGPTPro/comments/1i8793k/usage_of_knowledge_files_when_creating_a/

⁶ Introduction to Custom GPTs

<https://genai.byu.edu/introduction-to-custom-gpts>